

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

In [2]:

```
train_df=pd.read_csv(r"C:\Users\Gowthami\Downloads\Mobile_Price_Classification_train.csv")
train_df
```

Out[2]:

| | battery_power | blue | clock_speed | dual_sim | fc | four_g | int_memory | m_dep | mobile_wt | n_cores |
|------|---------------|------|-------------|----------|-----|--------|------------|-------|-----------|---------|
| 0 | 842 | 0 | 2.2 | 0 | 1 | 0 | 7 | 0.6 | 188 | |
| 1 | 1021 | 1 | 0.5 | 1 | 0 | 1 | 53 | 0.7 | 136 | |
| 2 | 563 | 1 | 0.5 | 1 | 2 | 1 | 41 | 0.9 | 145 | |
| 3 | 615 | 1 | 2.5 | 0 | 0 | 0 | 10 | 0.8 | 131 | |
| 4 | 1821 | 1 | 1.2 | 0 | 13 | 1 | 44 | 0.6 | 141 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1995 | 794 | 1 | 0.5 | 1 | 0 | 1 | 2 | 0.8 | 106 | |
| 1996 | 1965 | 1 | 2.6 | 1 | 0 | 0 | 39 | 0.2 | 187 | |
| 1997 | 1911 | 0 | 0.9 | 1 | 1 | 1 | 36 | 0.7 | 108 | |
| 1998 | 1512 | 0 | 0.9 | 0 | 4 | 1 | 46 | 0.1 | 145 | |
| 1999 | 510 | 1 | 2.0 | 1 | 5 | 1 | 45 | 0.9 | 168 | |

2000 rows × 11 columns



In [3]:

```
test_df=pd.read_csv(r"C:\Users\Gowthami\Downloads\Mobile_Price_Classification_test.csv")
test_df
```

Out[3]:

| | id | battery_power | blue | clock_speed | dual_sim | fc | four_g | int_memory | m_dep | mobile_wt |
|-----|------|---------------|------|-------------|----------|-----|--------|------------|-------|-----------|
| 0 | 1 | 1043 | 1 | 1.8 | 1 | 14 | 0 | 5 | 0.1 | 193 |
| 1 | 2 | 841 | 1 | 0.5 | 1 | 4 | 1 | 61 | 0.8 | 191 |
| 2 | 3 | 1807 | 1 | 2.8 | 0 | 1 | 0 | 27 | 0.9 | 186 |
| 3 | 4 | 1546 | 0 | 0.5 | 1 | 18 | 1 | 25 | 0.5 | 96 |
| 4 | 5 | 1434 | 0 | 1.4 | 0 | 11 | 1 | 49 | 0.5 | 108 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 995 | 996 | 1700 | 1 | 1.9 | 0 | 0 | 1 | 54 | 0.5 | 170 |
| 996 | 997 | 609 | 0 | 1.8 | 1 | 0 | 0 | 13 | 0.9 | 186 |
| 997 | 998 | 1185 | 0 | 1.4 | 0 | 1 | 1 | 8 | 0.5 | 80 |
| 998 | 999 | 1533 | 1 | 0.5 | 1 | 0 | 0 | 50 | 0.4 | 171 |
| 999 | 1000 | 1270 | 1 | 0.5 | 0 | 4 | 1 | 35 | 0.1 | 140 |

1000 rows × 21 columns

In [4]:

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   battery_power    2000 non-null   int64
1   blue             2000 non-null   int64
2   clock_speed      2000 non-null   float64
3   dual_sim         2000 non-null   int64
4   fc               2000 non-null   int64
5   four_g          2000 non-null   int64
6   int_memory       2000 non-null   int64
7   m_dep           2000 non-null   float64
8   mobile_wt       2000 non-null   int64
9   n_cores          2000 non-null   int64
10  pc               2000 non-null   int64
11  px_height        2000 non-null   int64
12  px_width         2000 non-null   int64
13  ram              2000 non-null   int64
14  sc_h             2000 non-null   int64
15  sc_w             2000 non-null   int64
16  talk_time        2000 non-null   int64
17  three_g          2000 non-null   int64
18  touch_screen     2000 non-null   int64
19  wifi             2000 non-null   int64
20  price_range      2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.3 KB
```

In [5]:

```
test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   id                  1000 non-null   int64  
 1   battery_power       1000 non-null   int64  
 2   blue                1000 non-null   int64  
 3   clock_speed         1000 non-null   float64 
 4   dual_sim            1000 non-null   int64  
 5   fc                  1000 non-null   int64  
 6   four_g              1000 non-null   int64  
 7   int_memory          1000 non-null   int64  
 8   m_dep               1000 non-null   float64 
 9   mobile_wt           1000 non-null   int64  
10   n_cores              1000 non-null   int64  
11   pc                   1000 non-null   int64  
12   px_height            1000 non-null   int64  
13   px_width             1000 non-null   int64  
14   ram                  1000 non-null   int64  
15   sc_h                 1000 non-null   int64  
16   sc_w                 1000 non-null   int64  
17   talk_time            1000 non-null   int64  
18   three_g              1000 non-null   int64  
19   touch_screen         1000 non-null   int64  
20   wifi                 1000 non-null   int64  
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

In [6]:

```
x=test_df.drop('dual_sim',axis=1)
y=test_df['dual_sim']
```

In [7]:

```
x=test_df.drop('dual_sim',axis=1)
y=test_df['dual_sim']
```

In [8]:

```
train_df['blue'].value_counts()
```

Out[8]:

```
blue
0    1010
1     990
Name: count, dtype: int64
```

In [9]:

```
test_df['blue'].value_counts()
```

Out[9]:

blue
1 516
0 484
Name: count, dtype: int64

In [10]:

```
T={"three_g":{"Yes":1,'No':0}}  
train_df=train_df.replace(T)  
print(train_df)
```

| | battery_power | blue | clock_speed | dual_sim | fc | four_g | int_memory | | |
|------|---------------|------|-------------|----------|----|--------|------------|---|--|
| 0 | 842 | 0 | 2.2 | 0 | 1 | 0 | 7 | \ | |
| 1 | 1021 | 1 | 0.5 | 1 | 0 | 1 | 53 | | |
| 2 | 563 | 1 | 0.5 | 1 | 2 | 1 | 41 | | |
| 3 | 615 | 1 | 2.5 | 0 | 0 | 0 | 10 | | |
| 4 | 1821 | 1 | 1.2 | 0 | 13 | 1 | 44 | | |
| ... | ... | ... | ... | ... | .. | ... | ... | | |
| 1995 | 794 | 1 | 0.5 | 1 | 0 | 1 | 2 | | |
| 1996 | 1965 | 1 | 2.6 | 1 | 0 | 0 | 39 | | |
| 1997 | 1911 | 0 | 0.9 | 1 | 1 | 1 | 36 | | |
| 1998 | 1512 | 0 | 0.9 | 0 | 4 | 1 | 46 | | |
| 1999 | 510 | 1 | 2.0 | 1 | 5 | 1 | 45 | | |

| | m_dep | mobile_wt | n_cores | ... | px_height | px_width | ram | sc_h | sc_w | |
|------|-------|-----------|---------|-----|-----------|----------|------|------|------|---|
| 0 | 0.6 | 188 | 2 | ... | 20 | 756 | 2549 | 9 | 7 | \ |
| 1 | 0.7 | 136 | 3 | ... | 905 | 1988 | 2631 | 17 | 3 | |
| 2 | 0.9 | 145 | 5 | ... | 1263 | 1716 | 2603 | 11 | 2 | |
| 3 | 0.8 | 131 | 6 | ... | 1216 | 1786 | 2769 | 16 | 8 | |
| 4 | 0.6 | 141 | 2 | ... | 1208 | 1212 | 1411 | 8 | 2 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 1995 | 0.8 | 106 | 6 | ... | 1222 | 1890 | 668 | 13 | 4 | |
| 1996 | 0.2 | 187 | 4 | ... | 915 | 1965 | 2032 | 11 | 10 | |
| 1997 | 0.7 | 108 | 8 | ... | 868 | 1632 | 3057 | 9 | 1 | |
| 1998 | 0.1 | 145 | 5 | ... | 336 | 670 | 869 | 18 | 10 | |
| 1999 | 0.9 | 168 | 6 | ... | 483 | 754 | 3919 | 19 | 4 | |

| | talk_time | three_g | touch_screen | wifi | price_range |
|------|-----------|---------|--------------|------|-------------|
| 0 | 19 | 0 | 0 | 1 | 1 |
| 1 | 7 | 1 | 1 | 0 | 2 |
| 2 | 9 | 1 | 1 | 0 | 2 |
| 3 | 11 | 1 | 0 | 0 | 2 |
| 4 | 15 | 1 | 1 | 0 | 1 |
| ... | ... | ... | ... | ... | ... |
| 1995 | 19 | 1 | 1 | 0 | 0 |
| 1996 | 16 | 1 | 1 | 1 | 2 |
| 1997 | 5 | 1 | 1 | 0 | 3 |
| 1998 | 19 | 1 | 1 | 1 | 0 |
| 1999 | 2 | 1 | 1 | 1 | 3 |

[2000 rows x 21 columns]

In [11]:

```
T={"three_g":{"Yes":1,'No':0}}
test_df=test_df.replace(T)
print(test_df)
```

```
..      ...      ...      ...      ..      ...      ...      ...      ...
995      0.5      170      ...      17      644      913      2121      14      8
996      0.9      186      ...      2      1152      1632      1933      8      1
997      0.5      80      ...      12      477      825      1223      5      0
998      0.4      171      ...      12      38      832      2509      15      11
999      0.1      140      ...      19      457      608      2828      9      2
```

```
      talk_time  three_g  touch_screen  wifi
0              2         0             1      0
1              7         1             0      0
2             10         0             1      1
3              7         1             1      0
4              7         1             0      1
..      ...      ...      ...      ...
995             15         1             1      0
996             19         0             1      1
997             14         1             0      0
998              6         0             1      0
999              3         1             0      1
```

51000 rows x 11 columns

In [12]:

```
x=train_df.drop('dual_sim',axis=1)
y=train_df['dual_sim']
```

In [13]:

```
x=test_df.drop('dual_sim',axis=1)
y=test_df['dual_sim']
```

In [14]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7,random_state=42)
x_train.shape,x_test.shape
```

Out[14]:

((700, 20), (300, 20))

In [15]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[15]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In [16]:

```
rf=RandomForestClassifier()
```


In [17]:

```
params={'max_depth':[2,3,5,10,20],  
        'min_samples_leaf':[5,10,20,50,100,200],  
        'n_estimators':[10,25,30,50,100,200]}
```

In [18]:

```
from sklearn.model_selection import GridSearchCV  
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')  
grid_search.fit(x_train,y_train)
```

Out[18]:



In [19]:

```
grid_search.best_score_
```

Out[19]:

```
0.54
```

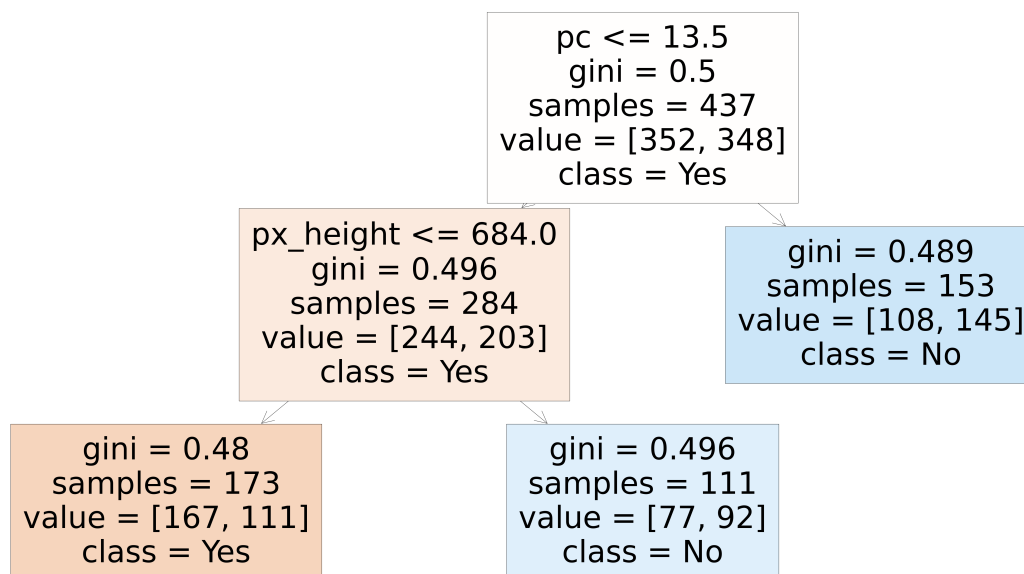
In [20]:

```
rf_best=grid_search.best_estimator_  
print(rf_best)
```

```
RandomForestClassifier(max_depth=20, min_samples_leaf=100, n_estimators=25)
```

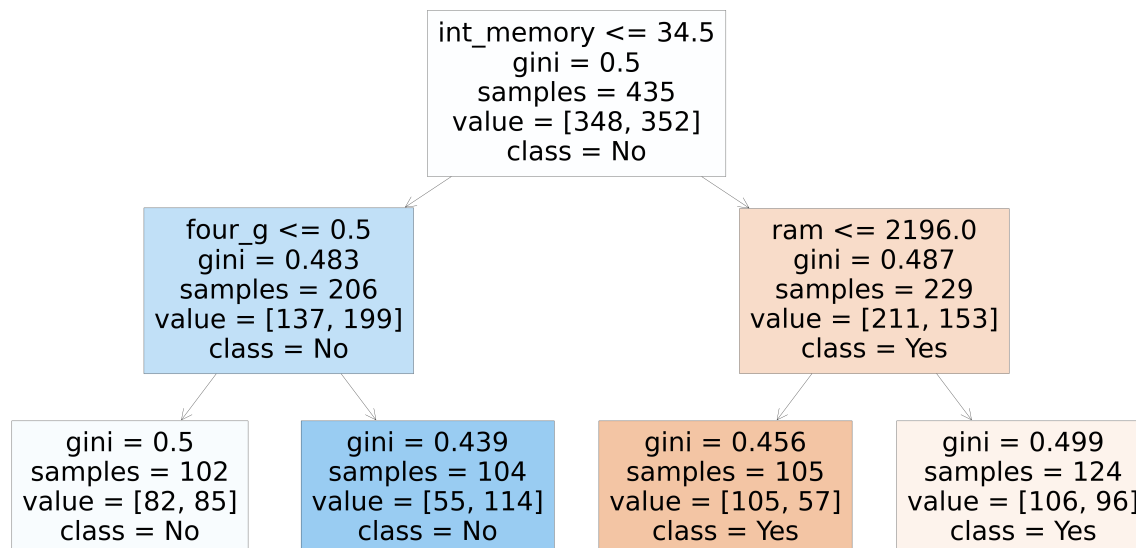
In [21]:

```
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],filled=True)
```



In [22]:

```
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=['Yes','No'],filled=True)
```



In [23]:

```
rf_best.feature_importances_
```

Out[23]:

```
array([0.05580668, 0.08252183, 0.          , 0.05483273, 0.0743369 ,
        0.02484387, 0.07418941, 0.04860052, 0.04644177, 0.01337144,
        0.12439139, 0.06679401, 0.03921233, 0.07888818, 0.01534584,
        0.10772848, 0.05103547, 0.          , 0.04165914, 0.          ])
```

In [24]:

```
imp_df=pd.DataFrame({'Varname':x_train.columns,"Imp":rf_best.feature_importances_})  
imp_df.sort_values(by="Imp",ascending=False)
```

Out[24]:

| | Varname | Imp |
|----|---------------|----------|
| 10 | pc | 0.124391 |
| 15 | sc_w | 0.107728 |
| 1 | battery_power | 0.082522 |
| 13 | ram | 0.078888 |
| 4 | fc | 0.074337 |
| 6 | int_memory | 0.074189 |
| 11 | px_height | 0.066794 |
| 0 | id | 0.055807 |
| 3 | clock_speed | 0.054833 |
| 16 | talk_time | 0.051035 |
| 7 | m_dep | 0.048601 |
| 8 | mobile_wt | 0.046442 |
| 18 | touch_screen | 0.041659 |
| 12 | px_width | 0.039212 |
| 5 | four_g | 0.024844 |
| 14 | sc_h | 0.015346 |
| 9 | n_cores | 0.013371 |
| 2 | blue | 0.000000 |
| 17 | three_g | 0.000000 |
| 19 | wifi | 0.000000 |

In []: